

# Python for ML

Week November 7th, 2020

Prepared by: Moe Fadae & Najem Bouazza

Exercise from last session

```
In [ ]: '''  
review today's lesson and answer the following question :  
  
Grab the number 34 from the this array:  
  
array = [23,4,[12,45,7,[2,67,8,4,8,[234,76,7,[231,6,6,[5,34,5]]]]]]  
  
'''
```

```
In [46]: import numpy as np
```

```
In [47]: array = np.array([23,4,[12,45,7,[2,67,8,4,8,[234,76,7,[231,6,6,[5,34,5]]]]]])
```

```
In [54]: array[2][3][5][3][3][1]
```

```
Out[54]: 34
```

```
In [ ]:
```

```
In [1]: import numpy as np  
array = np.array([23,4,[12,45,7,[2,67,8,4,8,[234,76,7,[231,6,6,[5,34,5]]]]]])
```

```
In [8]: array[2][3][5][3][3][1]
```

```
Out[8]: 34
```

Question from last session : How to delete an element from a list??

```
In [55]: L = [1,2,3,4,5]  
L.remove(L[3])  
L
```

```
Out[55]: [1, 2, 3, 5]
```

## I - NumPy

```
In [1]: import numpy as np
```

## - NumPy array Indexing / Slicing :

### Get an element in a list, 1d array, 2d array

```
In [19]: import numpy as np

my_list = [1, 'w', 98, "Get"]
array_1d = np.array([1,3,4,5,6,12])
array_2d = np.array([[1,2,3],[23,54,4],[21,54,67]])

print("Index for a list's element:",my_list[2])
print("Index for a 1d array's element:", array_1d[4])
print("Index for a 2d array's element:", array_2d[1][2])
```

```
Index for a list's element: 98
Index for a 1d array's element: 6
Index for a 2d array's element: 4
```

### Get a slice in a list, 1d array, 2d array

```
In [20]: print("slice from 1d array:", array_1d[1:])
print("slice from 2d array_:", array_2d[1:2,1:3])
```

```
slice from 1d array: [ 3  4  5  6 12]
slice from 2d array_: [[54  4]]
```

## II - Pandas :

```
In [2]: import pandas as pd
```

### 1 - Series and DataFrames

Series are similar to numpy arrays: what differentiate Pandas Series from Numpy arrays is that we can index Series using labels.

```
In [ ]:
```

```
In [60]: labels = ["a", "b", "c"]  
my_data = [11, "hello", 33]  
array = np.array(my_data)  
S = pd.Series(data= my_data)  
S
```

```
Out[60]: 0      11  
        1    hello  
        2      33  
        dtype: object
```

```
In [62]: S[2]
```

```
Out[62]: 33
```

```
In [63]: # Convert a NumPy array into a pandas Series  
  
pd.Series(array)
```

```
Out[63]: 0      11  
        1    hello  
        2      33  
        dtype: object
```

```
In [13]: #Convert a dictionary into Series  
  
d = {"one": 11, "two": 22, "three": 33}  
pd.Series(d)
```

```
Out[13]: one      11  
        two      22  
        three     33  
        dtype: int64
```

```
In [14]: # Series can hold any type of data (Strings / numbers /built in functions ...)  
  
pd.Series(data = [sum, print, len])
```

```
Out[14]: 0    <built-in function sum>  
        1    <built-in function print>  
        2    <built-in function len>  
        dtype: object
```

**Pandas** provides in-memory 2d table object called Dataframe

```
In [64]: from numpy.random import rand  
  
df = pd.DataFrame(data = rand(4,3), index = ["first", "second", "Third", "Fourth"]
```

In [65]: df

Out[65]:

	a	b	c
<b>first</b>	0.690696	0.772040	0.000844
<b>second</b>	0.948885	0.979848	0.785738
<b>Third</b>	0.644835	0.390349	0.462431
<b>Fourth</b>	0.293501	0.181962	0.328634

In [66]: df["b"]

Out[66]: first      0.772040  
second      0.979848  
Third      0.390349  
Fourth      0.181962  
Name: b, dtype: float64

In [ ]:

**A Dataframe is a bench of Series charing the same index !!!!**

In [41]: *#df[name of column] to grab a specific column*  
*# or df.columnname*  
df["b"]  
df.a *# it's not recomanded to use this method as it induces confusion with meth*  
  
*#df.(hit Tab on the keyboard) to see the available methods*

Out[41]: first      0.546217  
second      0.539157  
Third      0.742983  
Fourth      0.717028  
Name: a, dtype: float64

In [69]: df[["a", "b"]]

Out[69]:

	a	b
<b>first</b>	0.690696	0.772040
<b>second</b>	0.948885	0.979848
<b>Third</b>	0.644835	0.390349
<b>Fourth</b>	0.293501	0.181962

```
In [42]: # Grab mutiple columns, use the brakets and a list of columns names inside the brackets
df[["a", "c"]]
```

```
Out[42]:
```

	a	c
first	0.546217	0.177938
second	0.539157	0.126492
Third	0.742983	0.306025
Fourth	0.717028	0.361270

```
In [79]: df["new_column"] = df["a"]*100
df
```

```
Out[79]:
```

	a	b	c	Sum	new_column
first	0.690696	0.772040	0.000844	1.462736	69.069582
second	0.948885	0.979848	0.785738	1.928732	94.888462
Third	0.644835	0.390349	0.462431	1.035184	64.483483
Fourth	0.293501	0.181962	0.328634	0.475463	29.350100

```
In [80]: # Create new column :
#df["Sum"] = df["a"] + df["b"]
#df

# to eliminate a column : df.drop(column name, axis = 1)

df.drop("new_column", axis = 1)
```

```
Out[80]:
```

	a	b	c	Sum
first	0.690696	0.772040	0.000844	1.462736
second	0.948885	0.979848	0.785738	1.928732
Third	0.644835	0.390349	0.462431	1.035184
Fourth	0.293501	0.181962	0.328634	0.475463

```
In [81]: df
```

```
Out[81]:
```

	a	b	c	Sum	new_column
first	0.690696	0.772040	0.000844	1.462736	69.069582
second	0.948885	0.979848	0.785738	1.928732	94.888462
Third	0.644835	0.390349	0.462431	1.035184	64.483483
Fourth	0.293501	0.181962	0.328634	0.475463	29.350100

In [82]: *# Pandas wants you to confirm if you want to drop a column by using "inplace"*

```
df.drop("new_column", axis = 1, inplace = True)
df
```

Out[82]:

	a	b	c	Sum
<b>first</b>	0.690696	0.772040	0.000844	1.462736
<b>second</b>	0.948885	0.979848	0.785738	1.928732
<b>Third</b>	0.644835	0.390349	0.462431	1.035184
<b>Fourth</b>	0.293501	0.181962	0.328634	0.475463

In [60]: `df.drop("first", inplace = True)`

In [61]: `df`

Out[61]:

	a	b	c
<b>second</b>	0.539157	0.278046	0.126492
<b>Third</b>	0.742983	0.508975	0.306025
<b>Fourth</b>	0.717028	0.820784	0.361270

In [83]: *#Why we refer to axis = 0 for rows and axis = 1 for columns : because it's taken from*  
`df.shape`

Out[83]: (4, 4)

In [ ]: *##### Selecting Rows*  
*# we have to use a method to do this*  
*# 1 df.loc[label of index, rows] Label based index*  
*# 2 df.iloc[index number] numerical based index*

In [88]: `df.iloc[2]`  
`#df.loc["first"]`  
`# Error ? Why ?`

Out[88]: a 0.644835  
b 0.390349  
c 0.462431  
Sum 1.035184  
Name: Third, dtype: float64

In [63]: `df.loc["Third"]`

Out[63]: a 0.742983  
b 0.508975  
c 0.306025  
Name: Third, dtype: float64

```
In [64]: df.iloc[1]
```

```
Out[64]: a    0.742983
         b    0.508975
         c    0.306025
         Name: Third, dtype: float64
```

```
In [89]: df
```

```
Out[89]:
```

	a	b	c	Sum
<b>first</b>	0.690696	0.772040	0.000844	1.462736
<b>second</b>	0.948885	0.979848	0.785738	1.928732
<b>Third</b>	0.644835	0.390349	0.462431	1.035184
<b>Fourth</b>	0.293501	0.181962	0.328634	0.475463

```
In [90]: df.loc["second", "b"]
```

```
Out[90]: 0.9798478061198371
```

```
In [91]: df.iloc[1,1]
```

```
Out[91]: 0.9798478061198371
```

```
In [65]: # Selecting a subset of rows/columns
```

```
df.loc["Third", "b"]
```

```
Out[65]: 0.5089752424453037
```

```
In [68]: df.loc[["second", "Third"], ["a", "c"]]
```

```
Out[68]:
```

	a	c
<b>second</b>	0.539157	0.126492
<b>Third</b>	0.742983	0.306025

```
In [76]: # Conditional Selection
```

```
cond = df > 0.5
```

```
df[cond]
```

```
Out[76]:
```

	a	b	c
<b>second</b>	0.539157	NaN	NaN
<b>Third</b>	0.742983	0.508975	NaN
<b>Fourth</b>	0.717028	0.820784	NaN

```
In [96]: df[df["a"]>0.6]["b"]
```

```
Out[96]: first      0.772040  
second    0.979848  
Third     0.390349  
Name: b, dtype: float64
```

```
In [94]: df1 = df[df["a"]>0.6][["b","c"]]  
df1
```

```
Out[94]:
```

	b	c
Third	0.508975	0.306025
Fourth	0.820784	0.361270

```
In [ ]:
```

## III - Functions

### Create a function :

using the **def** keyword

It's a block of code which only runs when it is called

```
In [97]: def my_function():  
         print("Hello from a function")
```

### Calling a function :

use the function name followed by parenthesis

```
In [98]: my_function()  
  
Hello from a function
```

```
In [100]: def my_function(name = "default"):  
          print("Hello my name is {}".format(name))
```

```
In [102]: my_function("Jack")  
  
Hello my name is Jack
```

```
In [ ]:
```



```
In [104]: def Add(a, b):  
          """  
          This function returns the sum of 2 numbers  
          """  
          print("The first number is:", a)  
          print("The second number is:",b)  
          print("the sum of {} + {} is :".format(a,b), a+b)  
          return a + b
```

```
In [106]: Add(4,8)
```

```
The first number is: 4  
The second number is: 8  
the sum of 4 + 8 is : 12
```

```
Out[106]: 12
```

```
In [ ]:
```

## IV - Linear Regression :


### Practice on *house\_price* data

In this part, you will understand and learn how to implement the following Machine Learning Regression model:

#### Simple Linear Regression

- Fit a line to a dataset of observations
- Use this line to predict unobserved values

Exemple of Linear regression model :



## STEP 1 : Importing the libraries

```
In [107]: # Importing the Libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## STEP 2 : Importing the Dataset

```
In [110]: pd.read_csv("house_price.csv")
```

```
Out[110]:
```

	LotArea	SalePrice
0	8450	208500
1	9600	181500
2	11250	223500
3	9550	140000
4	14260	250000
...	...	...
1402	7917	175000
1403	13175	210000
1404	9042	266500
1405	9717	142125
1406	9937	147500

1407 rows × 2 columns

```
In [119]: # Importing the dataset
dataset = pd.read_csv('house_price.csv')
X = dataset.iloc[:, 0].values # .values convert dataframe into numpy array
y = dataset.iloc[:, 1].values # dataset.to_numpy()

dataset
```

```
Out[119]:
```

	LotArea	SalePrice
0	8450	208500
1	9600	181500
2	11250	223500
3	9550	140000
4	14260	250000
...	...	...
1402	7917	175000
1403	13175	210000
1404	9042	266500
1405	9717	142125
1406	9937	147500

1407 rows × 2 columns

### STEP 3 : Splitting the dataset into training\_set and test\_set

```
In [113]: # Splitting the dataset into the Training set and Test set  
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

## STEP 4 : Training the model using the training\_set

```
In [115]: # Training the Simple Linear Regression model on the Training set  
# import linear_model from sklearn  
  
from sklearn.linear_model import LinearRegression  
  
regressor = LinearRegression() # created our model  
  
regressor.fit(X_train, y_train) # Train the model
```

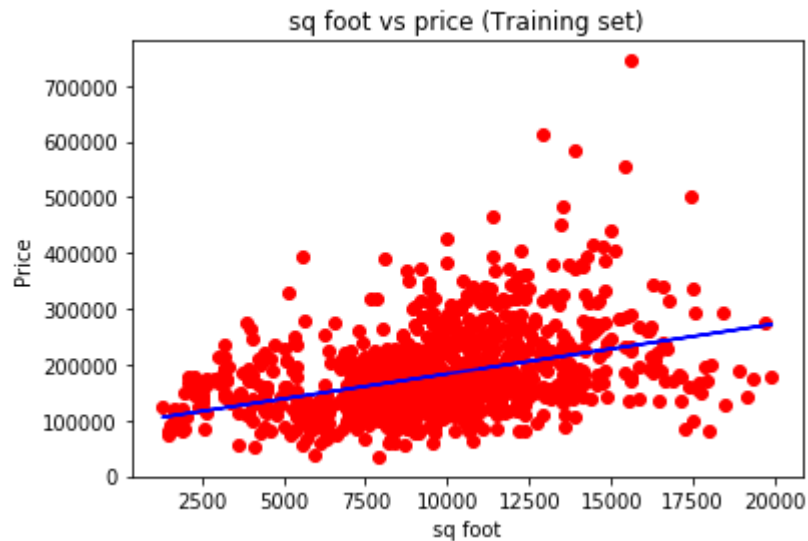
```
Out[115]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

## STEP 5 : Predicting the results for the test\_set

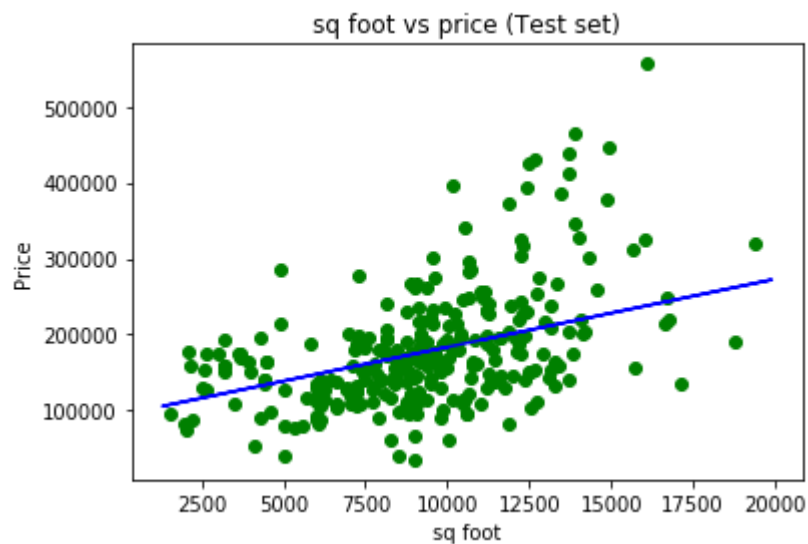
```
In [132]: # Predicting the Test set results  
y_pred = regressor.predict(X_test)  
#print(y_pred)
```

## STEP 6 : Visualizing the results

```
In [127]: # Visualising the Training set results
plt.scatter(X_train, y_train, color = 'red')
plt.plot(X_train, regressor.predict(X_train), color = 'blue')
plt.title('sq foot vs price (Training set)')
plt.xlabel('sq foot')
plt.ylabel('Price')
plt.show()
```



```
In [128]: # Visualising the Test set results
plt.scatter(X_test, y_test, color = 'green')
plt.plot(X_train, regressor.predict(X_train), color = 'blue')
plt.title('sq foot vs price (Test set)')
plt.xlabel('sq foot')
plt.ylabel('Price')
plt.show()
```



```
In [129]: # Model Intercept and Slope
intercept = regressor.intercept_
slope = regressor.coef_

print ("slope is: ", slope, " and intercept is: " , intercept)

slope is: [8.95119579] and intercept is: 93999.9307251301
```


```
In [131]: # R-squared
r_squared = regressor.score(X.reshape(-1,1), y.reshape(-1,1))
print ("R_squared is: ", r_squared)

R_squared is: 0.1746813314901382
```

```
In [32]: #Plot residuals
# calculate residuals
resi = y_test - y_pred

plt.plot (X_test, resi, 'bo')
```

Out[32]: [ <matplotlib.lines.Line2D at 0xc3c0748>]



```
In [134]: #Predict a value using the trained model

regressor.predict(np.array([[15000]]))
```

Out[134]: array([228267.86757217])

In [ ]: