

# Logistic Regression .2

Week November 28th, 2020

Prepared by: Moe Fadae & Najem Bouazza

## Introduction: ¶

In this Tutorial, we will train a logistic regression model to predict if customer is more likely to make a purchase.

The model will be trained based on 3 key features : Age, Gender, Salary.

This is a binary (binomial) classification : purchase did happen: 1, not purchased : 0

At the end, we will take a look into the results and the metrics used to evaluate the model, and how to manage the threshold

or the classification boundary.

The data has been retrieved from kaggle.

## I - Data Exploration & Preprocessing

```
In [1]: #Importing necessary packages  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [2]: #Reading dataset  
dataset = pd.read_csv("purchase.csv")
```

In [3]: dataset

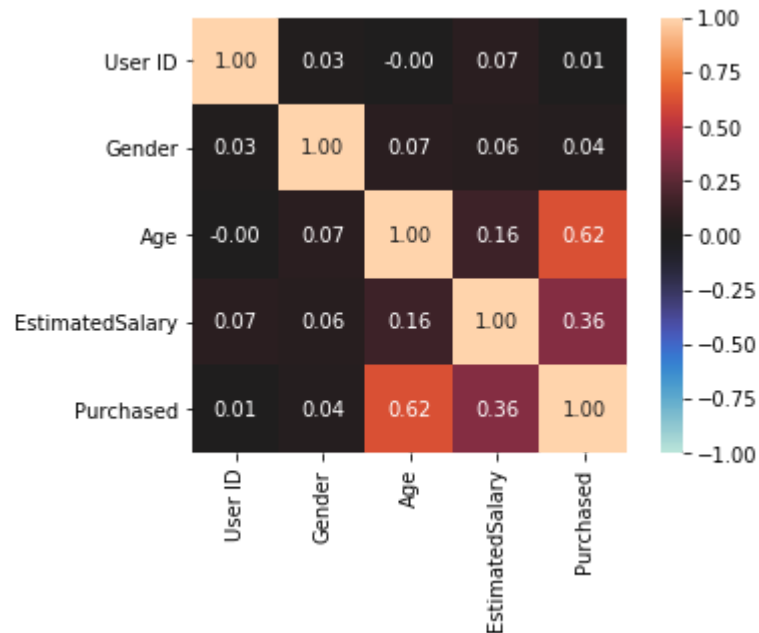
Out[3]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

In [4]: *# You can use NOMINAL class to view the correlation (including the categorical variables)*

```
from dython import nominal
nominal.associations(dataset, nominal_columns=['Gender'])
```



```
Out[4]: {'corr':
          User ID  Gender  Age  EstimatedSalary  Purcha
sed
User ID          1.000000  0.025249 -0.000721          0.071097  0.007120
Gender           0.025249  1.000000  0.073741          0.060435  0.042469
Age              -0.000721  0.073741  1.000000          0.155238  0.622454
EstimatedSalary  0.071097  0.060435  0.155238          1.000000  0.362083
Purchased        0.007120  0.042469  0.622454          0.362083  1.000000,
'ax': <matplotlib.axes._subplots.AxesSubplot at 0xb3a5b88>}
```

In [ ]:

In [5]: *#Checking for missing data*  
dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

In [6]: dataset.isna().any().sum()

Out[6]: 0

```
In [7]: #Properties of data
dataset.describe()
```

Out[7]:

	User ID	Age	EstimatedSalary	Purchased
<b>count</b>	4.000000e+02	400.000000	400.000000	400.000000
<b>mean</b>	1.569154e+07	37.655000	69742.500000	0.357500
<b>std</b>	7.165832e+04	10.482877	34096.960282	0.479864
<b>min</b>	1.556669e+07	18.000000	15000.000000	0.000000
<b>25%</b>	1.562676e+07	29.750000	43000.000000	0.000000
<b>50%</b>	1.569434e+07	37.000000	70000.000000	0.000000
<b>75%</b>	1.575036e+07	46.000000	88000.000000	1.000000
<b>max</b>	1.581524e+07	60.000000	150000.000000	1.000000

```
In [8]: #Encode the categorical feature - Gender using LabelEncoder
```


```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
dataset['Gender'] = le.fit_transform(dataset['Gender'])
dataset.head()
```

Out[8]:

	User ID	Gender	Age	EstimatedSalary	Purchased
<b>0</b>	15624510	1	19	19000	0
<b>1</b>	15810944	1	35	20000	0
<b>2</b>	15668575	0	26	43000	0
<b>3</b>	15603246	0	27	57000	0
<b>4</b>	15804002	1	19	76000	0

```
In [9]: #heatmap - correlation matrix of features
plt.figure(figsize=(10,8))
sns.set( font_scale= 2)
sns.heatmap(dataset.corr(),annot=True,cmap='inferno',mask=np.triu(dataset.corr()),
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0xb2b4b08>
```



```
In [10]: dataset.drop("User ID", axis =1, inplace =True)
```

```
In [11]: dataset
```

Out[11]:

	Gender	Age	EstimatedSalary	Purchased
0	1	19	19000	0
1	1	35	20000	0
2	0	26	43000	0
3	0	27	57000	0
4	1	19	76000	0
...	...	...	...	...
395	0	46	41000	1
396	1	51	23000	1
397	0	50	20000	1
398	1	36	33000	0
399	0	49	36000	1

400 rows × 4 columns

```
In [12]: #checking for duplicate samples  
dataset.duplicated().sum()
```

Out[12]: 20

```
In [13]: #dropping ALL duplicate values  
dataset.drop_duplicates(keep = False, inplace = True)
```

```
In [14]: dataset.shape
```


Out[14]: (362, 4)

```
In [15]: #Checking if any outliers
plt.figure(figsize=(10, 8))

plt.subplot(2,1,1)
sns.boxplot(dataset['Age'],color='yellow')


plt.subplot(2,1,2)
sns.boxplot(dataset['EstimatedSalary'], color='yellow')

plt.tight_layout()
```



```
In [16]: #Output distribution
plt.figure(figsize=(8, 6))
sns.countplot('Purchased', data=dataset)
```

Out[16]: <matplotlib.axes.\_subplots.AxesSubplot at 0xc9ba048>



**Define X and y, Convert them into arrays**

```
In [17]: X = dataset.drop("Purchased", axis = 1).values
y = dataset["Purchased"].values
```

In [ ]:

## II - Train & Evaluate the model

**Splitting the data**

```
In [18]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random
```

**Feature scaling**



```
In [19]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [20]: print(X_train[:3, :])

[[-0.96263527 -1.37394385  0.3233538 ]
 [-0.96263527  0.12620169  0.1205811 ]
 [-0.96263527  1.81386542 -1.29882782]]
```

### Create and fit the model

```
In [21]: from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression()

log_reg.fit(X_train, y_train)
```

```
Out[21]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)
```

### Predictions for X\_test

```
In [22]: y_pred = log_reg.predict(X_test)
y_pred_proba = log_reg.predict_proba(X_test)
```

```
In [23]: pd.DataFrame(y_pred_proba)
```

```
Out[23]:
```

	0	1
0	0.934922	0.065078
1	0.131239	0.868761
2	0.565517	0.434483
3	0.988452	0.011548
4	0.240856	0.759144
...	...	...
68	0.352628	0.647372
69	0.652208	0.347792
70	0.982579	0.017421
71	0.994265	0.005735
72	0.388475	0.611525

73 rows × 2 columns

### Changing the classification threshold

```
In [80]: def my_filter(x):
          if x > 0.9:
              return 1
          else:
              return 0
```

```
In [81]: y_pred_new = np.array([my_filter(x) for x in y_pred_proba[:,1]])
```


```
In [82]: y_pred_new
```

```
Out[82]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1,
                1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0])
```

### Confusion matrix for the default threshold:

```
In [83]: from sklearn.metrics import confusion_matrix
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cbar = False)
```


Out[83]: <matplotlib.axes.\_subplots.AxesSubplot at 0xfd86348>



```
In [84]: sns.set( font_scale=2)
labels =np.array(['True Neg','False Pos'],['False Neg','True Pos'])
sns.heatmap(confusion_matrix(y_test, y_pred), annot=labels,fmt='', cmap='Blues',

plt.xlabel('Predictions')
plt.ylabel('Actual Values')
```

Out[84]: Text(20.5, 0.5, 'Actual Values')



**Definitions :**

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 \times precision \times recall}{precision + recall}$$


$$accuracy = \frac{TP + TN}{TP + FN + TN + FP}$$

$$specificity = \frac{TN}{TN + FP}$$

**Confusion matrix for a new threshold:**

In [85]: `sns.heatmap(confusion_matrix(y_test, y_pred_new), annot=True, cbar = False)`

Out[85]: `<matplotlib.axes._subplots.AxesSubplot at 0xfb1e808>`



In [86]: `#Metrics based result`  
`from sklearn.metrics import classification_report`  
`print(classification_report(y_test, y_pred_new))`

	precision	recall	f1-score	support
0	0.72	0.90	0.80	48
1	0.62	0.32	0.42	25
accuracy			0.70	73
macro avg	0.67	0.61	0.61	73
weighted avg	0.68	0.70	0.67	73

```
In [87]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.75	0.79	48
1	0.60	0.72	0.65	25
accuracy			0.74	73
macro avg	0.72	0.73	0.72	73
weighted avg	0.76	0.74	0.74	73

```
In [88]: plt.figure(figsize = (16,7))
```

```
#Confusion matrix with labels
```

```
plt.subplot(1,3,1)
```

```
labels =np.array([[ 'TN', 'FP'],[ 'FN', 'TP']])
```

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=labels,fmt='', cmap='Blues',
```

```
plt.xlabel('Predictions')
```

```
plt.ylabel('Actual Values')
```

```
#Confusion matrix (0.5)
```

```
plt.subplot(1,3,2)
```

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cmap='Greens', cbar = f
```

```
plt.title("Default - 50%")
```


```
#Confusion matrix (new threshold)
```

```
plt.subplot(1,3,3)
```

```
sns.heatmap(confusion_matrix(y_test, y_pred_new), annot=True,cmap='Greens', cbar
```

```
plt.title("New threshold")
```

```
plt.tight_layout()
```



```
In [ ]:
```

In [92]: `dataset.head(3)`

Out[92]:

	Gender	Age	EstimatedSalary	Purchased
0	1	19	19000	0
1	1	35	20000	0
2	0	26	43000	0

In [ ]:

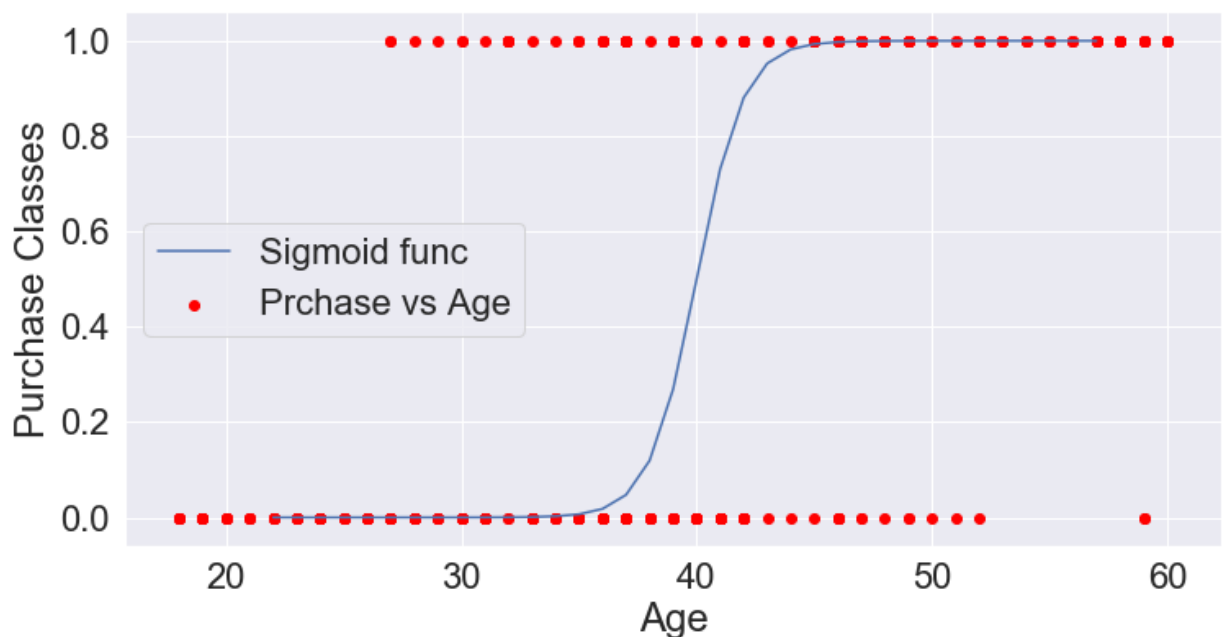
In [ ]:

### Plotting a sigmoid function and a scatter of data

```
In [119]: plt.figure(figsize = (12,6))
x = np.arange(-18,18)
plt.plot(x + 40, 1/(1+np.exp(-x)) , label = "Sigmoid func")
plt.scatter(x = dataset["Age"], y = dataset["Purchased"], color = "red" , label = "Prchase vs Age")

plt.xlabel("Age")
plt.ylabel("Purchase Classes")
plt.legend(loc = 6)
#dataset.plot.scatter(x = "EstimatedSalary", y = "Purchased")
```

Out[119]: <matplotlib.legend.Legend at 0x10576a48>



In [ ]:

